

INTRODUCTION TO RECOMMENDATION SYSTEMS WITH COLLABORATIVE FILTERING

AI FRAMEWORKS

Brendan Guillouet

6th December 2020

Institut National des Sciences Appliquées

INTRODUCTION

TABLE OF CONTENTS

Introduction

Neighborhood-based methods

Latent Vector-based methods

Factorization

Completion

Neural Network-based methods

TP

CONTEXT

- Customer relationship management.
- Appetence score.
 - Google double click, Criteo.
- Product recommendation

RECOMMENDATION SYSTEM

- Content Base. Using and Item metadata.
- Adaptive filtering (Bandit).
- Collaborative filtering.
 - Based on existing product X client relationship

2 types of solution.

- Based on neighborhood.
- Based on latent factor.
 - Matrices factorization.
 - Matrices completion
 - Neural Network

Difficulties

- How to evaluate the recommendation?
- Number of parameters to estimate.
 - Cold start, number of latent factors etc..

NEIGHBORHOOD-BASED METHODS

Two methods

1. User-User Filter

- Find a similarity metric between users.
- For a user u , find the set of k closest users S_u^k .
- Predict a rate, $\widehat{r}_{u,i}$ for an item i as a linear combination of known rates $r_{u',i}$ for $u' \in S_u^k$.

Two methods

1. User-User Filter

- Find a similarity metric between users.
- For a user u , find the set of k closest users S_u^k .
- Predict a rate, $\widehat{r}_{u,i}$ for an item i as a linear combination of known rates $r_{u',i}$ for $u' \in S_u^k$.

2. Item-Item Filter

- Find a similarity between items.
- For a user i , find the set of k closest users S_i^k .
- Predict a rate, $\widehat{r}_{u,i}$ for an user u as a linear combination of known rates $r_{u,i'}$ for $i' \in S_i^k$.

USER-USER FILTERS

Main assumption : customers with a similar profile will have similar tastes.

GENERIC FORMULA.

$$\hat{r}_{u,i} = \frac{\sum_{u' \in S_u^k} s(u, u') \cdot r_{u',i}}{\sum_{u' \in S_u^k} |s(u, u')|}$$

- $r_{u,i}$ is the rate given by user u to item i .
- s is a proximity score between u and u' .
- S_u^k is the set of k closest neighbors of user u

USER-USER FILTERS

Main assumption : customers with a similar profile will have similar tastes.

BASIC FORMULA WITH MEAN.

$$\hat{r}_{u,j} = \bar{r}_{u,.} + \frac{\sum_{u' \in S_u^k} s(u, u') \cdot (r_{u',j} - \bar{r}_{u',.})}{\sum_{u' \in S_u^k} |s(u, u')|}$$

- $r_{u,i}$ is the rate given by user u to item i .
- s is a proximity score between u and u' .
- S_u^k is the set of k closest neighbors of user u
- $\bar{r}_{u,.}$ is the mean rate given by user u

USER-USER FILTERS

Main assumption : customers with a similar profile will have similar tastes.

BASIC FORMULA WITH Z-SCORE.

$$\hat{r}_{u,j} = \bar{r}_{u,.} + \sigma_{u,.} \cdot \frac{\sum_{u' \in S_u^k} s(u, u') \cdot \frac{(r_{u',j} - \bar{r}_{u',.})}{\sigma_{u',.}}}{\sum_{u' \in S_u^k} |s(u, u')|}$$

- $r_{u,i}$ is the rate given by user u to item i .
- s is a proximity score between u and u' .
- S_u^k is the set of k closest neighbors of user u
- $\bar{r}_{u,.}$ is the mean rate given by user u
- $\sigma_{u,.}$ is the standard deviation of rate given by user u

PEARSON CORRELATION COEFFICIENT

$$s(u, v) = \frac{\sum_{i \in I_u \cap I_v} (r_{u,i} - \bar{r}_u) \cdot (r_{v,i} - \bar{r}_v)}{\sqrt{\sum_{i \in I_u \cap I_v} (r_{u,i} - \bar{r}_u)^2} \cdot \sqrt{\sum_{i \in I_u \cap I_v} (r_{v,i} - \bar{r}_v)^2}}$$

SPEARMAN RANK CORRELATION COEFFICIENT

$$s(u, v) = \frac{\sum_{i \in I_u \cap I_v} (\tilde{r}_{u,i} - \bar{\tilde{r}}_u) \cdot (\tilde{r}_{v,i} - \bar{\tilde{r}}_v)}{\sqrt{\sum_{i \in I_u \cap I_v} (\tilde{r}_{u,i} - \bar{\tilde{r}}_u)^2} \cdot \sqrt{\sum_{i \in I_u \cap I_v} (\tilde{r}_{v,i} - \bar{\tilde{r}}_v)^2}}$$

$\tilde{r}_{u,i}$ is the rank of product i in the preferences of user u (the higher the score, the smaller the rank).

COSINE : BASED ON VECTORIAL SPACE STRUCTURE

$$s(u, v) = \cos(u, v) = \frac{\langle u, v \rangle}{\|u\| \cdot \|v\|}$$

The unknown values have to be filled with zeros to compute this score.

Main assumption : the customers will prefer products that share a high similarity with those already well appreciated.

GENERIC FORMULA.

$$\hat{r}_{u,i} = \frac{\sum_{i' \in S_i^k} s(i, i') \cdot r_{u,i'}}{\sum_{i' \in S_i^k} |s(i, i')|}$$

- $r_{u,i}$ is the rate given by user u to item i .
- s is a proximity score between i and i' .
- S_i^k is the set of k closest neighbors of item i .

ITEM-ITEM FILTERS

Main assumption : the customers will prefer products that share a high similarity with those already well appreciated.

BASIC FORMULA WITH MEAN.

$$\hat{r}_{u,i} = \bar{r}_{.,i} + \frac{\sum_{i' \in S_i^k} s(i, i') \cdot (r_{u,i'} - \bar{r}_{.,i'})}{\sum_{i' \in S_i^k} |s(i, i')|}$$

- $r_{u,i}$ is the rate given by user u to item i .
- s is a proximity score between i and i' .
- S_i^k is the set of k closest neighbors of item i .
- $\bar{r}_{.,i}$ is the mean rate given to item i

ITEM-ITEM FILTERS

Main assumption : the customers will prefer products that share a high similarity with those already well appreciated.

BASIC FORMULA WITH Z-SCORE.

$$\hat{r}_{u,i} = \bar{r}_{.,i} + \sigma_{.,i} \cdot \frac{\sum_{i' \in S_i^k} s(i, i') \cdot \frac{(r_{u,i'} - \bar{r}_{.,i'})}{\sigma_{.,i'}}}{\sum_{i' \in S_i^k} |s(i, i')|}$$

- $r_{u,i}$ is the rate given by user u to item i .
- s is a proximity score between i and i' .
- S_i^k is the set of k closest neighbors of item i .
- $\bar{r}_{.,i}$ is the mean rate given to item i
- $\sigma_{.,i}$ is the standard deviation of rate given to item i

PROXIMITY SCORES

Same proximity score can be used :

PEARSON CORRELATION COEFFICIENT

SPEARMAN RANK CORRELATION COEFFICIENT

COSINE : BASED ON VECTORIAL SPACE STRUCTURE

Another approach relies on the Bayesian paradigm.

CONDITIONAL PROBABILITY

$$s(i, j) = P[j \in B | iB]$$

Where B is the purchase history.

- Easy to interpret results.
- Choice of k can be done with cross-validation.
- Good performance on small dataset..
- ... but suffer from scalability problem as user base grows.
- Use with **SURPRISE** python library.

LATENT VECTOR-BASED METHODS

LATENT VECTOR-BASED METHODS

FACTORIZATION

MATRIX FACTORIZATION

Let $X \in \mathcal{M}_{N \times M}$ be the matrix of user/item notations where N is the number of users and M the number of items.

MAIN IDEA : Approximate X as a product of two matrices $W \in \mathcal{M}_{N \times K}$ and $H \in \mathcal{M}_{K \times M}$

items


	1		3			5			5		4	
			5	4			4			2	1	3
users	2	4		1	2		3		4	3	5	
		2	4		5			4			2	
			4	3	4	2					2	5
	1		3		3			2			4	

X

MATRIX FACTORIZATION

Let $X \in \mathcal{M}_{N \times M}$ be the matrix of user/item notations where N is the number of users and M the number of items.


MAIN IDEA : Approximate X as a product of two matrices $W \in \mathcal{M}_{N \times K}$ and $H \in \mathcal{M}_{K \times M}$



MATRIX FACTORIZATION

Let $X \in \mathcal{M}_{N \times M}$ be the matrix of user/item notations where N is the number of users and M the number of items.

MAIN IDEA : Approximate X as a product of two matrices $W \in \mathcal{M}_{N \times K}$ and $H \in \mathcal{M}_{K \times M}$




Each entry of X can be approximated with the following formula :

$$X_{i,j} = W_{i,\cdot}^T \cdot H_{\cdot,j} = \sum_{k \in K} w_{i,k} \cdot h_{k,j}$$

INTERPRETATION

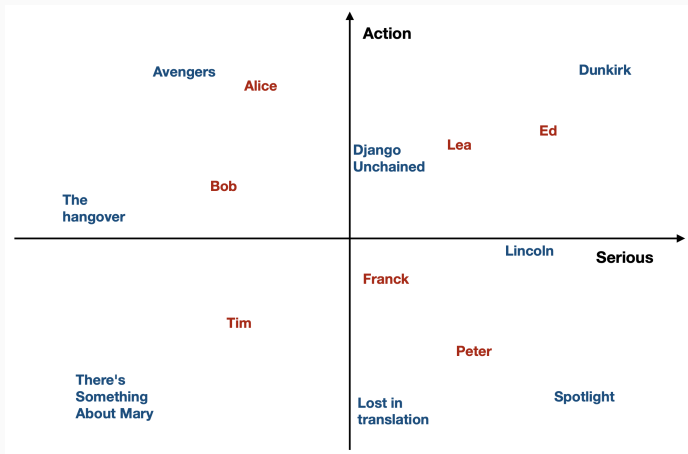
The K latent vectors can be interpreted as embedding.

Example : Movie recommendation



Latent factors can be interpreted as the genre of the movie.

INTERPRETATION



Similar users/movies get embedded nearby in the embedding space . It's possible to compute similarities in this space.

To recover W and H the following minimization problem is introduced :

$$\min_{W, H \geq 0, rk(W)=r, rk(H)=r} L(X, WH) + \lambda_w P(W) + \lambda_h P(H)$$

- L is the loss function (measure accuracy of prediction)
- λ_w, λ_h are penalization parameter
- P is the penalty function.

NON NEGATIVE MATRIX FACTORIZATION

To recover W and H the following minimization problem is introduced :

$$\min_{W, H \geq 0, rk(W)=r, rk(H)=r} L(X, WH) + \lambda_w P(W) + \lambda_h P(H)$$

- L is the loss function (measure accuracy of prediction)
- λ_w, λ_h are penalization parameter
- P is the penalty function.

Constraint of positivity on W and $H \implies$ **Non Negative Matrix Factorization**.

NMF algorithms generally solve problem iteratively.

LEE AND SEUNG [1999] Euclidean loss, no regularization.

$$H_{k,j} \leftarrow H_{k,j} \frac{(W^T X)_{k,j}}{(W^T W H)_{k,j}}, \quad W_{i,k} \leftarrow W_{i,k} \frac{(X H^T)_{i,k}}{(V H H^T)_{i,k}},$$

BRUNET ET AL. [2004] Kullback-Leibler loss, no regularization.

$$H_{k,j} \leftarrow H_{k,j} \frac{\sum_l \frac{W_{l,k} X_{l,j}}{(W H)_{l,j}}}{\sum_l W_{l,k}}, \quad W_{i,k} \leftarrow W_{i,k} \frac{\sum_l \frac{H_{k,l} X_{i,l}}{(W H)_{i,l}}}{\sum_l H_{k,l}}$$

Most algorithm available in R NMF package are variation of those algorithms.

Example with L_2 regularization and euclidean loss function.

$$\min_{W, H \geq 0, rk(W)=r, rk(H)=r} \sum_{i,j} (x_{i,j} - w_i^T h_j)^2 + \sum_{i=1}^N \lambda_w ||w_i||^2 + \sum_{j=1}^M \lambda_h ||h_j||^2$$

A non convex problem.

Example with L_2 regularization and euclidean loss function.

$$\min_{W, H \geq 0, rk(W)=r, rk(H)=r} \sum_{i,j} (x_{i,j} - w_i^T h_j)^2 + \sum_{i=1}^N \lambda_w ||w_i||^2 + \sum_{j=1}^M \lambda_h ||h_j||^2$$

A non convex problem.

SOLUTION : Alternate Least Square.

$\forall i$, fix all variables except w_i and solve :

$$\operatorname{argmin}_{w_i} \sum_j (x_{i,j} - w_i^T h_j)^2 + \lambda_w ||w_i||^2$$

Do the same with fix h_j , $\forall j$.

It is easy to show that :

$$\operatorname{argmin}_{w_i} \sum_j (x_{i,j} - w_i^T h_j)^2 + \lambda_w ||w_i||^2$$

Has the solution

$$w_i = (\sum_j h_j h_j^T + \lambda_w I_k)^{-1} (\sum_j x_{i,j} v_j)$$

This is very similar to regularized least squares regression.

These algorithms has been widely studied and a lot of different implementation exist.

- **NMF** R package (9 implementation among Kim and Park [2007], Lee and Seung [1999], Brunet et al. [2004]).
 - get a lot of interpretation tools such (consensus matrix, dendogram etc..)
- **MLlib** SPARK (Koren et al. [2009]).
 - Implementation that handle missing value.
- **Surprise** PYTHON (Zhang et al. [2006], Luo et al. [2014]).
- **Scikit-learn** PYTHON (Lee and Seung [1999], Hoyer [2004])

- SVD decomposition can also provide a solutions for recomendation.

$$X = UDV^*$$

- Good property as unity and optimal solution are guaranteed for fixed rank.
- Optimize the rank.
- NMF minimization problem does not always lead to well posed problem.
- Cold start problem not discussed but it's a major research point.
- Most of the implementation of these algorithm treat missing rates as zero.

LATENT VECTOR-BASED METHODS

COMPLETION

Mazumder et al. [2010] re-write the minimization problem as

$$\min_M \frac{1}{2} \|P_\Omega(X) - P_\Omega(M)\|_F^2 + \lambda \|M\|_*$$

where

- Ω contain the pairs of indices (i, j) where X is observed
- $P_\Omega(X)_{i,j} = X_{i,j}$ if $X_{i,j}$ is known, 0 otherwise.
- $\|M\|_*$ is the nuclear norm of M

If \hat{M} solves problem, then it satisfies the following stationarity condition :

$$\hat{M} = S_{\lambda}(Z)$$

where

$$Z = P_{\Omega}(X) + P_{\Omega^{\perp}}(\hat{M})$$

and

$$S_{\lambda}(Z) = UD_{\lambda}V', \quad D_{\lambda} = \text{diag}[(d_1 - \lambda)_+, \dots, (d_r - \lambda)_+]$$

Reconstruct $S_{\lambda}(Z) = UD_{\lambda}V^T$ from SVD decomposition is called “soft-thresholded SVD”.

For sufficiently large λ , D_{λ} will be rank-reduced, and hence so will be $UD_{\lambda}V^T$.

Pseudo Code

Initialize \hat{M}

Iterate over the following steps :

- Compute $Z = P_{\Omega}(X) + P_{\Omega^{\perp}}(\hat{M})$.
- Compute $\hat{M} = S_{\lambda}(Z)$.

- We only use know value of X .
- Much more faster than NMF.
- implementation within the **softImpute** R package.

LATENT VECTOR-BASED METHODS

NEURAL NETWORK-BASED METHODS

As for NMF the main idea is minimize the objective function

$$\sum_{i,j \in \Omega} (r_{i,j} - w_i^T h_j)_2^2 + \lambda(||W||_2^2 + ||H||_2^2)$$

where

- the embedding w_i of a movie i in an embedding space of size k .
- the embedding h_j of an item j in an embedding space of size k .

As for NMF the main idea is minimize the objective function

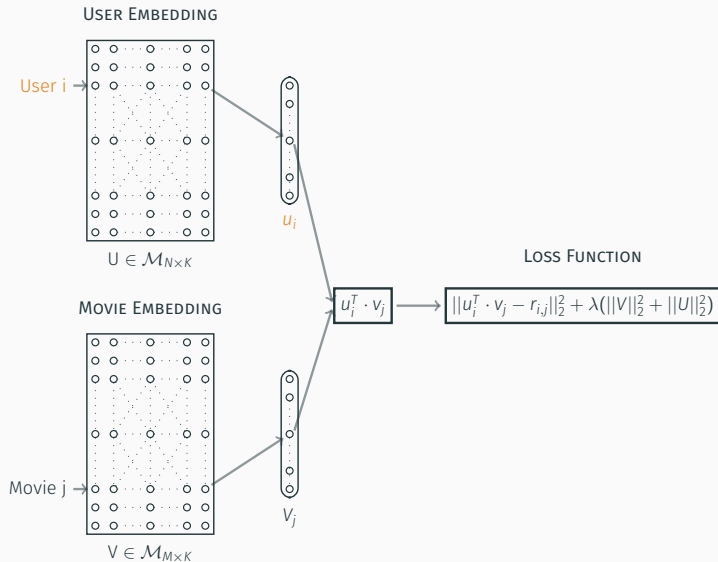
$$\sum_{i,j \in \Omega} (r_{i,j} - w_i^T h_j)_2^2 + \lambda(||W||_2^2 + ||H||_2^2)$$

where


- the embedding w_i of a movie i in an embedding space of size k .
- the embedding h_j of an item j in an embedding space of size k .

SOLUTION Use neural network architecture and framework to solve it!

ARCHITECTURE



DEEP ARCHITECTURE



```
import tensorflow.keras.layers as kl
user_embedding = kl.Embedding(output_dim=embedding_size,
                              input_dim=input_dim,
                              input_length=1)
```

- embedding_size = K the size of latent vector
- input_dim = M The size of the user
- input_length = 1 (i.e the indice).

All weight of *user_embedding* are trainable and tuned during backprob.

- Handle missing data (as it just optimize over existing user item couple
- Good performance comparing to other methods (especially if you have GPU).
- Hard to interpret the function you designed to compare the vector...
- ...but you can still interpret the embedding vectors!

TP

We will use the **MOVIELENS** dataset to compare the different methods.

- It's produced by the groupLens company (<https://grouplens.org/>).
- Various dataset :
 - Stable one. 25 Millions of rates. Use for challenge and benchmark.
 - Small one. 100K rates. Use for education or research. (Not stable).
- Contains various metadata that we won't use in this TP.
 - Genre, Tag, Age of user etc...

1. Neighborhood-based methods

- *1-Python-Neighborhood-MovieLens.ipynb*
 - Discover the data and surprise library. Apply Neighborhood based data.

2. Latent Vector-based methods

- *2-R-Facto-MovieLens.ipynb*
 - NMF package on toy dataset. SoftImpute on movielens.
- *2-Pyspark-Facto-MovieLens* (OPTIONAL)
- *3-Python-Neural-MovieLens*
 - Apply Neural recommender system on movielens dataset. Visualize latent vector.

<http://wikistat.fr/pdf/st-m-datSc3-colFil.pdf>

https://cse.iitk.ac.in/users/piyush/courses/ml_autumn16/771A_lec14_slides.pdf

https://m2dsupsdclass.github.io/lectures-labs/slides/03_recommender_systems/index.html

RÉFÉRENCES

- Jean-Philippe Brunet, Pablo Tamayo, Todd R Golub, and Jill P Mesirov. Metagenes and molecular pattern discovery using matrix factorization. *Proceedings of the national academy of sciences*, 101(12) :4164–4169, 2004.
- Patrik O Hoyer. Non-negative matrix factorization with sparseness constraints. *Journal of machine learning research*, 5(Nov) :1457–1469, 2004.
- Hyunsoo Kim and Haesun Park. Sparse non-negative matrix factorizations via alternating non-negativity-constrained least squares for microarray data analysis. *Bioinformatics*, 23(12) :1495–1502, 2007.

- Yehuda Koren, Robert Bell, and Chris Volinsky. Matrix factorization techniques for recommender systems. *Computer*, (8) :30–37, 2009.
- Daniel D Lee and H Sebastian Seung. Learning the parts of objects by non-negative matrix factorization. *Nature*, 401(6755) :788, 1999.
- Xin Luo, Mengchu Zhou, Yunni Xia, and Qingsheng Zhu. An efficient non-negative matrix-factorization-based approach to collaborative filtering for recommender systems. *IEEE Transactions on Industrial Informatics*, 10(2) :1273–1284, 2014.
- Rahul Mazumder, Trevor Hastie, and Robert Tibshirani. Spectral regularization algorithms for learning large incomplete matrices. *Journal of machine learning research*, 11(Aug) :2287–2322, 2010.

Sheng Zhang, Weihong Wang, James Ford, and Fillia Makedon. Learning from incomplete ratings using non-negative matrix factorization. In *Proceedings of the 2006 SIAM international conference on data mining*, pages 549–553. SIAM, 2006.